

Samee Arif

asamee@umich.edu | sameearif.com | GoogleScholar | GitHub | LinkedIn

EDUCATION

Ph.D. in Computer Science and Engineering

University of Michigan - Ann Arbor (UMich)

Advised by Professor Rada Mihalcea

Aug. 2025 - May 2030

GPA: 4.00/4.00

Bachelor of Science in Computer Science

Lahore University of Management Sciences (LUMS)

Sep. 2019 - May 2023

RESEARCH & TECHNICAL EXPERIENCE

Graduate Research Assistant | Language and Information Technologies (UMich)

Jan. 2026 - Present

- Conducting research on LLM safety, including jailbreaks, child-facing AI, cultural alignment, and multi-agent reliability.

Research Associate | Center for Speech and Language Technologies (LUMS)

June 2024 - Jan. 2025

- Proposed a novel multi-agent framework for automated preference optimization dataset generation leading to an 11% improvement in preference alignment.
- Built Kahaani, a multimodal multi-agent storytelling system that generates child-friendly stories using LLMs, TTS, TTV, and TTM grounded in Freytag's Pyramid and Propp's narrative functions, and evaluated through human studies.

GenAI and Software Engineer | ActualAIz

Aug. 2023 - June 2024

- Engineered an AI-driven multimodal assistant that streamlined the university selection process and statement of purpose writing, resulting in 100+ placements in top universities worldwide.

Research Assistant | LUMS

Aug. 2021 - May 2023

- Conducted evaluation of language models on Urdu NLP benchmarks, identifying gaps in high-quality Urdu datasets.
- Proposed EATS for improved answer alignment and constructed an Urdu QA corpus comprising 124,745 QA pairs.
- Developed the first Urdu Conversational Speech dataset for ASR benchmarking; fine-tuned and evaluated ASR models.
- Created the first Urdu Newspaper OCR dataset and fine-tuned language models for text recognition.

Teaching Assistant

- **Computer Science I: Fundamentals** (University of Iowa)
- **Machine Learning** (LUMS)
- **Computational Problem Solving** (LUMS)

Jan. 2025 - May 2025

Sep. 2022 - Dec. 2022

Sep. 2021 - Dec. 2021

PROJECTS

Incremental Completion Attacks on Aligned Language Models

- Proposed a jailbreak framework that elicits harmful outputs by generating one word at a time toward a harmful objective, gradually shifting the model's internal safety state and bypassing alignment through context accumulation.
- Achieved 87.31% ASR on Llama-3-8B and 78.01% ASR on Gemma-3-12B on AdvBench.

Developmentally-Aware Evaluation of Child-Facing LLMs

- Developed KIDBench, a developmental-psychology-grounded benchmark evaluating 13 LLMs across single- and multi-turn child interactions, revealing 10–30% gains from explicit age cues and up to 24% multi-turn degradation.
- Trained KIDGuardLlama, achieving 0.87 QWK against frontier-LLM judgments, and KIDLlama on 22K+ child-safe responses using SFT and Critique-GRPO; human evaluators preferred KIDLlama over the strongest baseline.

Hierarchical Cultural Reasoning & Value Alignment in LLMs

- Developing a Cultural Tree framework that structures country-level norms and values into hierarchical representations to analyze cultural reasoning patterns and value alignment in large language models.
- Conducting cross-cultural evaluation to measure value consistency, reasoning trace alignment, and generalization across countries, using normalized norm/value taxonomies derived from Cultural Atlas data.

RESEARCH GRANTS

Thinking Machines Research Grant

Aug 2026

VESSL AI Compute Grant

May 2026

OpenAI Research Access Program

July 2024

OpenAI Research Access Program

May 2024

PUBLICATIONS

- [1]. Samee Arif, et al., [One Word at a Time: Incremental Completion Decomposition Breaks LLM Safety](#). COLM 2026.
- [2]. Samee Arif, et al., [The Age of Curiosity Meets the Age of AI: Benchmarking Child Safety in LLMs](#). Preprint 2026.
- [3]. Samee Arif, et al., [With a Grain of SALT: Are LLMs Fair Across Social Dimensions?](#). TrustNLP at ACL 2026
- [4]. Samee Arif, et al., [From Press to Pixels: Evolving Urdu Text Recognition](#). LREC 2026.
- [5]. Samee Arif, et al., [Kahaani: A Multimodal Co-Creative Storytelling System](#). SRW at EACL 2026.
- [6]. Samee Arif, et al., [The Fellowship of the LLMs: Multi-LLM Synthetic Preference Optimization](#). GEM2 at ACL 2025.
- [7]. Samee Arif, et al., [WER We Stand: Benchmarking Urdu ASR Models](#). COLING 2025.
- [8]. Samee Arif, et al., [Generalists vs. Specialists: Evaluating Large Language Models for Urdu](#). EMNLP 2024.
- [9]. Samee Arif, et al., [UQA: Corpus for Urdu Question Answering](#). LREC-COLING 2024.

TECHNICAL SKILLS

Technical Stack: PyTorch, Hugging Face, vLLM, CUDA, LLM Fine-tuning, RLHF/GRPO, Multi-Agent Systems, Git, Docker